

Data Structures & Algorithms for Problem Solving (CS1.304)


Lecture # 16 : MST

Avinash Sharma

Center for Visual Information Technology (CVIT),
IIIT Hyderabad


Spanning Trees

- We will now consider another famous problem in graphs.
- Imagine providing connectivity to a set of cities.
- Each highway connects two cities
- In reality, each highway requires a certain cost to be built.



Spanning Trees

- So, there is a trade-off here.
- How to provide connectivity while minimizing the total cost of building the highways.
- The weights on the edges indicate the cost of building that highway.
- The total cost of connectivity = sum of all the built up highway.
- Minimize this cost.



Spanning Trees

- Formalize the problem as follows.
 - Let $G = (V, E, W)$ be a weighted graph.
 - Find a subgraph G' of G that is **connected** and has the **smallest cost**
 - Cost is defined as the sum of the edge weights of edges in G' .
-

Spanning Trees

- Observation I : If G' has a cycle and is connected, then there exists a G'' , which is also a subgraph of G and is connected so that
 - $\text{cost}(G'') < \text{cost}(G')$
 - To get G'' , simply break at least one cycle of G' .
 - Hence, the optimal G' shall have no cycles and is connected.
 - Suggests that G' is a tree.
-

Spanning Trees

- Two keywords : spanning and tree.
 - Some notation: A subgraph G' of G is called a **spanning subgraph** if $V(G') = V(G)$.
 - A spanning subgraph G' of G that is also a tree is called as a **spanning tree** of G .
-

Spanning Trees

- Consider the problem: Find a spanning tree of G that has the least cost.
 - Such a spanning tree is also called as a **minimum cost spanning tree** of G . Often one refers to this as the minimum spanning tree, or MST for short.
-

MST

- Let us now think of devising an algorithm to construct an MST of a given weighted graph G .
 - There are several approaches, but let us consider a bottom-up approach.
 - Let us start with a graph (tree) that has no edges and add edges successively.
 - Every new edge we add should not create a cycle.
 - Further, the total cost of the final tree should be the least possible.
-

MST

- Suggests that we should prefer edges of smaller weight.
 - But should not add edges that create cycles.
- Indeed, that is intuitive and turns out that is correct too.
 - we will skip the proof of this.



MST Algorithm

Algorithm MST(G)

begin

sort the edges of G in increasing order of weight as e_1, e_2, \dots, e_m

$k = 1$; $V(T) = V(G)$; $E(T) = \Phi$

while $|E(T)| < n-1$ do

if **$E(T) \cup e_k$ does not have a cycle** then

$E(T) = E(T) \cup e_k$


end-if

$k = k + 1$;

end-while




End.

MST Practice Problem



MST

- List of edges by weight
 - bc, ab, ac, cg, cd, de, bf, af, ad, ef, dg



MST Algorithm Analysis

- The algorithm we devised is called the Kruskal's algorithm.
- Belongs to a class of algorithms called greedy algorithms.
- How do we analyze our algorithm?
 - Need to know how to implement the cycle checker.



MST Algorithm Analysis

- How quickly can we find if a given graph has a cycle?
 - $O(m+n)$ is possible using DFS.
 - Notice that if the graph is a forest, then $m = O(n)$.
 - So, can be done in $O(n)$ time.
 - Also, need to try all m edges in the worst case.
 - So the time required in this case is $O(mn)$.
-

MST Algorithm Analysis

- Too high in general.
 - But, advanced data structures exist to bring the time down very close to $O(m+n)$.
 - Cannot be covered in this class.
 - We will show an approach that takes us almost there.
-

Advanced Data Structures

- An abstract problem:
 - Given n elements, grouped into a collection of disjoint sets S_1, S_2, \dots, S_k , design a data structure to:
 - Find the set to which an element belongs
 - Combine two sets
 - The abstract problem finds applications in several settings:
 - Spanning tree algorithm of Kruskal
 - Graph connected components
 - Least common ancestors
 - ...
-

Notations for Disjoint Sets

- Imagine a collection $S = \{S_1, S_2, \dots, S_k\}$ of sets.
 - Each set has a **representative** element
 - Some member of the set, typically.
 - Depending on application, can be
 - The smallest numbered element
 - A number ...
 - Typical operations
 - **MakeSet(x)** - Creates new set whose only member is x. The representative is x
 - **Union(x, y)** - Unites set S_x containing x and set S_y containing y into a new set S and removes S_x and S_y from the collection.
 - **FindSet(x)** - Returns representative of the set holding x
-


Some Notations

- Two parameters
 - n : The number of MakeSet operations.
 - m : The total number of MakeSet, Union, and Find operations.
 - Some observations
 - Each Union operation reduces the number of sets by 1.
 - When starting with n elements, at most $n-1$ Union operations.
 - Also, $m \geq n$.
 - Assume that the n MakeSet operations are the first n operations.
-

How to Implement the Operations?

- Option 1 : Use linked lists.
 - For every set, there is a linked list.
 - The representative of a set is the head of the list.
 - Every element also stores a pointer to the representative.
 - There is a tail pointer indicating where to append.
-


Example



Operations

- MakeSet(x): Create a new linked list.
 - FindSet(x) : Can be answered via the direct pointer
 - Union(x, y) : Can append the list of x to the list of y.
 - But have to update the pointer for each element in the list of x.
-


Adjacency List Example



Application to Connected Components


- Problem: Given an undirected graph $G = (V, E)$, partition V into disjoint sets V_1, V_2, \dots, V_k , so that two vertices u and v are in the same partition if and only if there is a path between u and v .
 - Several ways to solve this problem
 - This may not be the best way!
 - Example follows.
-

Example




- Algorithm:
 - For each vertex v
 - $\text{MakeSet}(v)$
 - For each edge vw
 - $\text{Union}(v,w)$




Example



Example



Example



Operations

- How difficult is it to append the lists?
 - Claim: There exists a sequence of m operations on n objects so that the total time required for the entire sequence of operations is $O(n^2)$.
 - After the first n MakeSet operations, call $\text{Union}(x_1, x_2)$, $\text{Union}(x_2, x_3)$, $\text{Union}(x_3, x_4)$, ..., $\text{Union}(x_{n-1}, x_n)$.
 - The k th union call takes time proportional to k .
 - Total time is therefore $O(n^2)$.
 - The average time per operation is also $O(n)$.
-

Application to Kruskal's Algorithm

- An average time of $O(n)$ is not helpful for Kruskal's algorithm.
- We have several Union calls and several FindSet calls.



Better Solution

- Most of the time spent is in the Union operation.
 - Can we modify the operation slightly?
 - Intuitively, it is easier to append a smaller list to a larger list.
 - Requires fewer updates.
 - Will the overall time decrease?
 - We will show that indeed it does.
-

The Weighted Union Heuristic

- Maintain the length of each list. Corresponds to the size of the set.
 - To perform $\text{Union}(x, y)$:
 - Append the list of x to the list of y if $\text{len}(x) < \text{len}(y)$
 - Append the list of y to the list of x otherwise.
 - A single Union operation can still take lot of time.
 - Union of two large lists, say of size $n/10$ each.
 - But, a sequence of operations may be not so expensive.
 - Hopefully.
-

Analysis

- How many times can an element change its representative?
 - Consider any element x .
 - If in an Union operation, the representative of x changes, then x is in the smaller list.
 - Why?
 - The first time this happens, the resulting list has at least 2 elements.
 - Next time, the resulting list has at least 4 elements.
-

Analysis

- In general, if the representative of x changes k times, then the resulting list has size at least 2^k .
 - The largest set can have a size of n .
 - Therefore, the representative of x cannot change more than $\log n$ times, over all the Union operations.
 - This applies to every element.
 - Therefore, over all Union operations, the total time spent is $O(n \log n)$.
-

Analysis

- Now, consider a sequence of m operations.
 - MakeSet and Find are $O(1)$ time operations.
 - Therefore, the total time is $O(m + n \log n)$.
 - The average time per operation is $O(\log n)$.
-

Application to Kruskal's Algorithm

- How does the above apply?



Application to Kruskal's Algorithm

- Do n MakeSet operations indicating that each vertex is in its own tree/set.
 - To check if $e = uv$ creates a cycle, check if $\text{FindSet}(u) = \text{FindSet}(v)$.
 - If not, add e to the current tree. Perform $\text{Union}(u, v)$ to merge the trees of u and v .
 - There are at most m FindSet operations.
 - Overall time is therefore bound by $O(m+n\log n)$.
-

MST – Another Approach

- The previous approach has to check for cycles every iteration.
 - Another approach that has a smaller runtime even with basic data structures.
 - Largely simplifies the solution.
-

MST – Another Approach

- The current approach is characterized by having a single tree T at any time.
 - In each iteration, T is extended by adding one vertex v not in T and one edge from v to some vertex in T .
 - Starting from a tree of one node, this process is repeated $n-1$ times.
-

MST – Another Approach

- Two questions:
 - How to pick the new vertex v ?
 - How to pick the edge to be added from v to some other vertex in T ?



MST – Another Approach

- The answers are provided by the following claims.
 - Claim 1: Let $G = (V, E, W)$ be a weighted undirected graph. Let v be **any** vertex in G . Let vw be the edge of **smallest weight** amongst all edges with one endpoint as v . Then vw is always contained in **any** MST of G .
-


MST – Another Approach

- Claim 1 can be shown in the following way.
- For each vertex v in G , there must be at least one edge in any MST.
- Considering the edge of the smallest weight is useful as it can decrease the cost of the spanning tree.




Generalizing Claim 1

- Let T be a subtree of some MST of an undirected weighted graph G .
- Consider edges uv in G such that u is in T and v is not in T .
- Of all such edges, let $e = xy$ be the edge with the smallest weight.
- Then $T \cup \{e\}$ is also a subtree of some MST of G .



Generalizing Claim 1

- Claim 2 allows us to expand a given sub-MST T .
- We can use Claim 2 to expand the current tree T .
- How to Start ?



Towards an Algorithm

- Let v be any vertex in the graph G . Pick v as the starting vertex to be added to T .
 - T now contains one vertex and no edges.
 - T is a subtree of some MST of G .
 - Now, apply Claim 2 and extend T .
-

Towards an Algorithm

Algorithm MST(G, v)

Begin

 Add v to T ;

 While T has less than $n - 1$ edges do

$w =$ vertex s.t. vw has the smallest weight
 amongst edges with one endpoint in T and
 another not in T .

 Add vw to T .

 End

End

Towards an Algorithm

- How to find w in the algorithm?
 - Need to maintain the weight of edges that satisfy the criteria.
 - A better approach:
 - Associate a key to every vertex
 - $\text{key}[v]$ is the smallest weight of edges with v as one endpoint and another in the current tree T .
 - $\text{key}[v]$ changes only when some vertex is added to T .
 - Vertex with the smallest $\text{key}[v]$ is the one to be added to T .
-

Towards an Algorithm

- Suggests that $\text{key}[v]$ need to be updated only when a new vertex is added to T .
 - Further, not all $\text{key}[v]$ may change in every iteration.
 - Only the neighbors of the vertex added to T .
 - Similar to Dijkstra's algorithm.
-

Towards an Algorithm

- Therefore, can maintain a heap of vertices with their $key[]$ values.
 - Initially, $key[v] = \text{infinity}$ for every vertex except the start vertex for which key value can be 0.
 - Perform DeleteMin on the heap. Let v be the result.
 - Update the $key[]$ value for neighbors w of v as:
 - $key[w] = \min\{key[w], W(vw)\}$
-

Algorithm using a Heap

Algorithm MST(G, u)

begin

for each vertex v do $\text{key}[v] = \text{infty}$.

$\text{key}[u] = 0$;

Add all vertices to a heap H .

While T has less than $n-1$ edges do

$v = \text{deleteMin}()$;

Add v to T via uv s.t. u is in T

For each neighbor w of v do

if $W(vw) > \text{key}[w]$ then $\text{DecreaseKey}(w)$

end

end

end

Algorithm using a Heap

- The algorithm is called as Prim's algorithm.
 - Runtime easy to analyze;
 - Each vertex deleted once from the heap. Each DeleteMin() takes $O(\log n)$ time. So, this accounts for a time of $O(n \log n)$.
 - Each edge may result in one call to DecreaseKey(). Over m edges, this accounts for a time of $O(m \log n)$.
 - Total time = $O((n+m) \log n)$.
-

Thank You

