

Data Structures & Algorithms for Problem Solving (CS1.304)

Sorting & Parallelism

Avinash Sharma

Center for Visual Information Technology (CVIT),
IIIT Hyderabad

Introduction

- Sorting is a fundamental concept in Computer Science.
 - several applications and a lot of literature.
 - We shall see two algorithms for sorting today
 - Try to introduce parallelism in sorting



QuickSort

- The quick sort algorithm designed by Tony Hoare is a simple yet highly efficient algorithm.
 - It works as follows:
 - Start with the given array A of n elements.
 - Consider a pivot, say $A[n]$.
 - Now, partition the elements of A into two arrays A_L and A_R such that:
 - the elements in A_L are less than $A[n]$
 - the elements in A_R are greater than $A[n]$.
 - Sort A_L and A_R , recursively.
-

QuickSort



- How to partition ?
 - Suppose we take each element, compare it with $A[n]$ and then move it to A_L or A_R accordingly.
 - Works in $O(n)$ time.
 - Can write the program easily.
 - But, recall that space is also an resource. The above approach requires extra space for the arrays A_L and A_R
 - A better approach exists.
-

QuickSort

```
Procedure Partition(A,n)
begin
  pivot = A[n];
  less = 0; more = 1;
  for more = 1 to n-1 do
    if A[more] < pivot then
      less++;
      swap(A[more], A[less]);
    end
  end
  swap (A[less+1], A[n]);
end
```


15	24	37	12	18	11	42	25	22
----	----	----	----	----	----	----	----	----

QuickSort



QuickSort

- Graphical Visualization of recursive partitioning,



Graphics Courtesy: Wikipedia

Analyzing Quick Sort

- Suppose we run quick sort with $A[n]$ as the pivot.
 - Let A_L and A_R be the two subarrays obtained after partitioning.
 - What is the time taken by quicksort?
 - As a recurrence relation, $T(n) = T(|A_L|) + T(|A_R|) + O(n)$.
 - To be able to solve this recurrence relation, need to know the sizes of arrays A_L and A_R .
-


Analyzing Quick Sort

- We know that $|A_L| + |A_R| = n-1$.
- But, if the pivot is such that all elements are smaller (or larger) than the pivot, then $|A_L|$ (or $|A_R|$) = $n-1$.
- The recurrence relation in that case is

$$T(n) = T(n-1) + O(n).$$

- Suppose the same situation happens over every recursive call. So, the above recurrence relation holds during every recursive call.
 - When will this happen ?
-

Analyzing Quick Sort



Analyzing Quick Sort

- Is it always that bad?
 - What if the pivot is such that each recursive iteration, the sizes of $|A_L|$ and $|A_R|$ is exactly the same?
 - The recurrence relation then stands as:
$$T(n) = 2T(n/2) + O(n).$$
 - Which element as the pivot ensures that the sizes of $|A_L|$ and $|A_R|$ are exactly the same?
 - Can that happen in every run when you pick a pivot to be the last element? Or the first element? Or even uniformly at random?
-

Analyzing Quick Sort

- In general, if the sizes of $|A_L|$ and $|A_R|$ are such that they are a constant away from each other, then the recurrence relation is:

$$T(n) = T(an) + T((1-a)n) + O(n)$$

where a is a constant < 1 .

- In practice, it turns out that most often the partitions are not too skewed.
 - So, quick sort runs in $O(n \log n)$ time almost always.
-

Merge Sort

- Another sorting technique.
- Based on the divide and conquer principle.
- We will first explain the principle and then apply it to merge sort.



Divide and Conquer

- Divide the problem P into $k \geq 2$ sub-problems P_1, P_2, \dots, P_k .
 - Solve the sub-problems P_1, P_2, \dots, P_k .
 - Combine the solutions of the sub-problems to arrive at a solution to P .
-

Divide and Conquer

- A useful paradigm with several applications.
- Examples include merge sort, convex hull, median finding, matrix multiplication, and others.
- Typically, the sub-problems are solved recursively.
 - Recurrence relation

$$T(n) = D(n) + \sum_i T(n_i) + C(n)$$

Divide time


Recursive cost

Combine time



Divide and Conquer

- Combination Procedure: Merge



Algorithm Merge

```
Algorithm Merge(L, R)
// L and R are two sorted arrays of size n each.
// The output is written to an array A of size 2n.
int i=1, j=1;
L[n+1] = R[n+1] = MAXINT; // so that index does not
                           // fall over
for k = 1 to 2n do
    if L[i] < R[j] then
        A[k] = L[i]; i++;
    else
        A[k] = R[j]; j++;
end-for
```

Time complexity is $O(n)$.


From Merging to Sorting

- How to use merging to finally sort?
 - Using the divide and conquer principle
 - Divide the input array into two halves.
 - Sort each of them.
 - Merge the two sub-arrays. This is indeed procedure Merge.
 - The algorithm can now be given as follows.
-

Algorithm Merge Sort

```
Algorithm MergeSort(A)
begin
    mid = n/2; //divide step
    L = MergeSort(A[1..mid]);
    R = MergeSort(A[mid+1..n]);
    Merge(L, R); //combine step
end-Algorithm
```

Divide & Conquer Merge Sort



- Example via merge sort: 1) Divide is split into two parts
2) Recursively solve each subproblem
-

Analyzing Merge Sort

- Recurrence relation for merge sort as:

$$T(n) = 2T(n/2) + O(n).$$

- This can be explained by the $O(n)$ time for merge and
 - The two subproblems obtained during the divide step each take $T(n/2)$ time.
 - Now use the general format for divide and conquer based algorithms.
 - Solving this recurrence relation is done using say the substitution method giving us $T(n) = O(n \log n)$.
 - Look at previous examples.
-

Introducing Parallelism in Computing

- My laptop has 8 cores.
 - What does it mean?
 - In principle, each core can run some instructions on its own independently.
 - So, while one core is possibly running the browser, the one can run an editor, the third can run a PDF reader, and the fourth can run a mail client, etc.
 - Is that helpful in all situations?
-

Introducing Parallelism in Computing

- There are two reasons why the previous model is not enough.
 - There are applications that are time critical and can use ALL the cores for themselves.
 - Examples include weather forecasting, cyber-security, and the like.
 - Secondly, while the number of cores is increasing, per core performance is dropping.
-

Introducing Parallelism in Computing

- So, how can MY program use ALL the cores simultaneously?
 - Can the architecture help?
 - Can the OS help?
 - Can the compiler help?
 - Can I help?
 - Very much. That is what parallel computing is all about.
 - How can I help?
 - Algorithm Designer
 - Programmer
 - Computer Science student,
 - All of the above.
-

Maxima of n Numbers


- A sequential program resembles the code below.

```
Program Maxima(A)
Begin
    int max = A[1];
    for i = 1 to n do
        if max < A[i]
            max = A[i]
        End-if
    End-for
End
```


- How do we run this program in parallel?
-

Maxima of n Numbers

- Suppose we compare pairs of numbers and record the maximum.
- The overall maximum is one of these local maxima.
- So, we could compare pairs of local maxima, to get more local maxima.
- Continue until only one pair remains with the property that the maxima of the input is the maximum of the pair that remains.



Maxima of n Numbers



Maxima of n Numbers

- Can you write the above idea as a parallel algorithm?
 - Some hints. Let n be the size of the array.
 - How many pairs are compared in the first iteration?
 - In the second iteration?
 - In the i th iteration?
 - How many iterations have to be executed to get the final maxima?
 - Which parts of the algorithm are parallel?
 - Within iterations or across iterations?
-

Parallelism in Computing

- Think of the sequential computer as a machine that executes jobs or instructions.
 - With more than one processor, can execute more than one job (instruction) at the same time.
 - Cannot however execute instructions that are dependent on each other.
 - This opens up a new world where computations have to be specified in parallel.
 - Sometimes have to rethink on known sequential approaches.
-

Parallelism in Computing

- How many independent operations can be done at a time?
 - Depends on the number of processors available.
 - Assume that as many as n , or $n/2$, processors are available.
 - Most often, our analysis suggests that the computation takes only $O(\log n)$ time, but we need n processors for this.
 - Cannot ensure that the number of processors also grow with the input size.
 - In practice, the number of processors on a machine does not change!
-

Parallelism in Computing

- The idea of the parallel algorithm is to show the extent of parallelism available in the computation.
- Plus, if there are fewer processors than what is required, can always simulate more processors.
- For instance, if there are p processors and n processors are required, then each of the p processors simulates the actions of n/p processors.



Parallel Merge Sort

- An algorithm is a sequence of tasks T_1, T_2, \dots
 - These tasks may have inter-dependencies,
 - Such as task T_i should be completed before task T_j for some i, j .
 - However, it is often the case that there are several algorithms where many tasks are independent of each other.
 - In some cases, the algorithm or the computation has to be expressed in that independent-task fashion.
-

Parallel Merge Sort

- In such a setting, one can imagine that tasks that are independent of each other can be done simultaneously, or in parallel.
- Let us think of arriving at a parallel merge sort algorithm.



Parallel Merge Sort

- What are the independent tasks in merge sort?
 - Sorting the two parts of the array.
 - This further breaks down to sorting four parts of the array, etc.
 - Eventually, every element of the array is a sorted sub-array.
 - So the main work is in merge itself.
-

Parallel Merge Sort

- So, we just have to figure out a way to merge in parallel.
- Recall the merge algorithm as we developed it earlier.
 - Too many dependent tasks.
 - Not feasible in a parallel model.

```
for k = 1 to 2n do
    if L[i] < R[j] then
        A[k] = L[i]; i++;
    else A[k] = R[j]; j++;
end-for
```

Parallel Merge Sort

- Need to rethink on a parallel merge algorithm
- Start from the beginning.
 - We have two sorted arrays L and R.
 - Need to merge them into a single sorted array A.



Parallel Merge Sort

- Need to rethink on a parallel merge algorithm
 - Start from the beginning.
 - We have two sorted arrays L and R.
 - Need to merge them into a single sorted array A.
 - Define the **rank** of an element x in a sorted array A as the number of elements of A that are smaller than x.
 - To merge L and R, need to know the rank of every element from L and R in the merged array $L \cup R$.
-

Parallel Merge Sort

- Importantly, for any x in L or R ,
 $\text{Rank}(x, L \cup R) = \text{Rank}(x, L) + \text{Rank}(x, R)$.
 - So, merging is equivalent to finding the two ranks on the right hand side.
-

Parallel Merge Sort

- Now, consider an element x in L at index k .
 - How many elements of L are smaller than x ?
 - $k-1$.
 - How many elements of R are smaller than x ?
 - No easy answer, but
 - can do binary search for x in R and get the answer.
 - Say k' elements in R are smaller than x .
-

Parallel Merge Sort

- How many elements in LUR are smaller than x ?
 - Precisely $k + k' - 1$.
 - So, in the merged output, what index should x be placed in?
 - precisely at $k+k'$.
 - Can this be done for every x in L ?
 - Yes, it is an independent operation.
 - Can this be done for every x in R also?
 - Yes, replace the roles of L and R .
 - All these operations are independent.
-

Parallel Merge Sort

L = [8 10 12 27]

R = [15 17 24 32]

Element	8	10	12	27	15	17	24	32
Rank in L	0	1	2	3	3	3	3	4
Rank in R	0	0	0	3	0	1	2	3
Rank in L U R	0	1	2	6	3	4	5	7

L U R = [8 10 12 15 17 24 27 32]

Parallel Merge Sort

```
Algorithm ParallelMergeSort(A)
```

```
Begin
```

```
    mid =  $n/2$ ; //divide step
```

```
    L = MergeSort(A[1..mid]);
```

```
    R = MergeSort(A[mid+1..n]);
```

```
    ParallelMerge(L, R); //combine step
```

```
end-Algorithm
```

Thank You

